

I'm gonna start with the Regularized linear models as my benchmark estimators! Models to be implemented in this notebook:

- Ridge Regression
- Lasso Regression
- Elastic

But first lets import some dependencies and our data

```
In [ ]: # importing the dependencies
import warnings
warnings.filterwarnings('ignore')

import matplotlib.pyplot as plt
%config InlineBackend.figure_format = 'retina'
plt.style.use('ggplot')
import pandas as pd
import numpy as np

import seaborn as sns
sns.set()

from sklearn.model_selection import train_test_split, KFold, learning_curve
from sklearn.model_selection import GridSearchCV, cross_val_score

from sklearn.linear_model import Ridge, Lasso, ElasticNet

from sklearn.preprocessing import PowerTransformer

from sklearn.metrics import mean_absolute_error, r2_score

from joblib import load,dump
```

```
In [ ]: # loading our data
df = pd.read_csv("./Data/data.csv",sep=",")
df.drop(['Unnamed: 0'], axis=1, inplace=True) # There were some formatting issues
# writing the csv
```

```
In [ ]: df.head()
```

```
Out[ ]:   DISTRICT  UPAZILA  STATION_ID  STATION_NAME  DATE  RAIN_FALL(mm)  LATITUDE  LONGITUDE
0  Bandarban      Lama     CL317        Lama  01-jan-2017       0.0      21.81      92.1
1  Bandarban      Lama     CL317        Lama  02-jan-2017       0.0      21.81      92.1
2  Bandarban      Lama     CL317        Lama  03-jan-2017       0.0      21.81      92.1
```

	DISTRICT	UPAZILA	STATION_ID	STATION_NAME	DATE	RAIN_FALL(mm)	LATITUDE	LONGITUDI
3	Bandarban	Lama	CL317	Lama	04-jan-2017	0.0	21.81	92.1
4	Bandarban	Lama	CL317	Lama	05-jan-2017	0.0	21.81	92.1

Lets make our X and y respectively

```
In [ ]: X = df['RAIN_FALL(mm)'].values.reshape(-1,1) # input feature
y = df['WATER_LEVEL(m)'].values.reshape(-1,1) # target feature
```

```
In [ ]: X.shape, y.shape
```

```
Out[ ]: ((1826, 1), (1826, 1))
```

Making train-test split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(
    X,y, test_size=0.2, random_state=17, shuffle=True
)
```

### Power Transform (Box-Cox)

```
In [ ]: pt = PowerTransformer(method='box-cox')
X_train_transformed = pt.fit_transform(X_train+0.000001)
X_test_transformed = pt.transform(X_test+0.000001)
```

```
In [ ]: pt.lambdas_
```

```
Out[ ]: array([-0.10662485])
```

```
In [ ]: y_train_transformed = pt.fit_transform(y_train+0.000001)
y_test_transformed = pt.transform(y_test+0.000001)
```

```
In [ ]: pt.lambdas_
```

```
Out[ ]: array([-5.34754941])
```

## Model Building

### Ridge Regression

*Linear least squares with l2 regularization.*

lets now initialize our CV and model with default parameters provided by sklearn!

```
In [ ]: kf10d = KFold(n_splits=5, shuffle=True, random_state=17)
ridge_regression = Ridge(random_state=17)
```

lets check our CV scores with default parameters

```
In [ ]: results = cross_val_score(
    ridge_regression,
    X_train_transformed,
    y_train_transformed,
    cv=kf10d,
    scoring='neg_mean_absolute_error'
)
results.mean()
```

```
Out[ ]: 0.605630841351809
```

Lets check for train test accuracy now

```
In [ ]: ridge_regression.fit(X_train_transformed,y_train_transformed)
```

```
Out[ ]: Ridge(random_state=17)
```

```
In [ ]: # accuracy on the train set
ridge_pred = ridge_regression.predict(X_train_transformed)
mean_absolute_error(y_train_transformed,ridge_pred)
```

```
Out[ ]: 0.6055762485546179
```

```
In [ ]: # accuracy on the test set
ridge_pred = ridge_regression.predict(X_test_transformed)
mean_absolute_error(y_test_transformed,ridge_pred)
```

```
Out[ ]: 0.6011885371212594
```

## Learning Curve

```
In [ ]: # Helper function
alphas = np.logspace(-2, 0, 20)
def plot_with_err(x, data, **kwargs):
    mu, std = data.mean(1), data.std(1)
    lines = plt.plot(x, mu, "-", **kwargs)
    plt.fill_between(
        x,
        mu - std,
        mu + std,
        edgecolor="none",
        facecolor=lines[0].get_color(),
        alpha=0.2,
    )
```

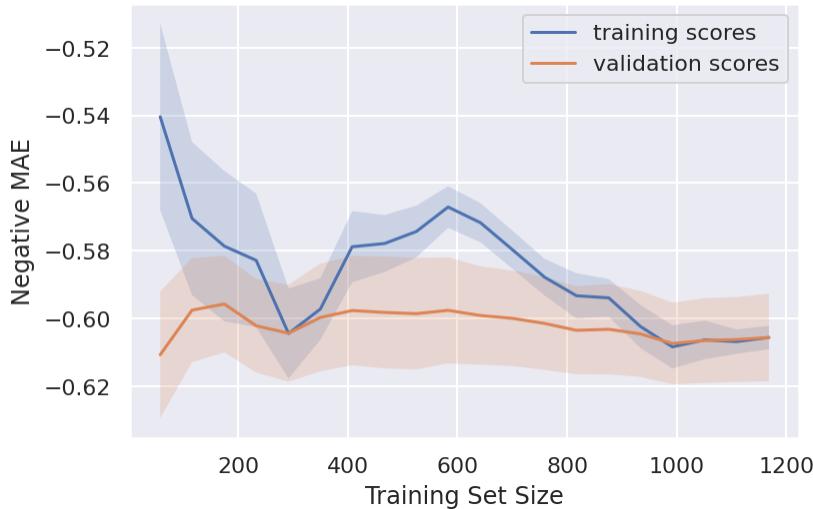
```
In [ ]: def plot_learning_curve():
```

```

train_sizes = np.linspace(0.05,1,20)
N_train, val_train, val_test = learning_curve(
    ridge_regression,X_train_transformed,y_train_transformed, train_sizes=train_sizes)
)
plot_with_err(N_train, val_train, label="training scores")
plot_with_err(N_train, val_test, label="validation scores")
plt.xlabel("Training Set Size")
plt.ylabel("Negative MAE")
plt.legend()
plt.grid(True);

```

In [ ]: plot\_learning\_curve()



Note:

- Theres seems like no variance problem, thus tuning the hyperparameter alpha wont do much good here! So lets move onto LASSO!
- by MAE 0.6 means that, on average model is 0.6 meter wrong on predicting the correct value of water level!

## Saving the Model

In [ ]: dump(ridge\_regression,'./Saved Model/RidgeRegression.joblib')

Out[ ]: ['./Saved Model/RidgeRegression.joblib']

## Lasso Regression

*Linear Model trained with L1 prior as regularizer (aka the Lasso).*

In [ ]: kfold = KFold(n\_splits=5,shuffle=True, random\_state=17)
lasso\_regression = Lasso(random\_state=17, max\_iter=10000)

In [ ]: results = cross\_val\_score(
 lasso\_regression,
 X\_train\_transformed,

```
        y_train_transformed,  
        cv=kfold,  
        scoring='neg_mean_absolute_error'  
)  
-results.mean()
```

Out[ ]: 0.8732801748385459

In [ ]: lasso\_regression.fit(X\_train\_transformed,y\_train\_transformed)

Out[ ]: Lasso(max\_iter=10000, random\_state=17)

In [ ]:  
*# accuracy on the train set*  
lasso\_pred = lasso\_regression.predict(X\_train\_transformed)  
mean\_absolute\_error(y\_train\_transformed,lasso\_pred)

Out[ ]: 0.8729811451213172

In [ ]:  
*# accuracy on the test set*  
lasso\_pred = lasso\_regression.predict(X\_test\_transformed)  
mean\_absolute\_error(y\_test\_transformed,lasso\_pred)

Out[ ]: 0.8666496419855901

## Learning Curve

In [ ]:

```
def plot_learning_curve():  
    train_sizes = np.linspace(0.05,1,20)  
    N_train, val_train, val_test = learning_curve(  
        lasso_regression,X_train_transformed,y_train_transformed, train_sizes=train_sizes  
)  
    plot_with_err(N_train, val_train, label="training scores")  
    plot_with_err(N_train, val_test, label="validation scores")  
    plt.xlabel("Training Set Size")  
    plt.ylabel("Negative MAE")  
    plt.legend()  
    plt.grid(True);
```

In [ ]: plot\_learning\_curve()



## Saving The Model

```
In [ ]: dump(lasso_regression,'./Saved Model/LassoRegression.joblib')
```

```
Out[ ]: ['./Saved Model/LassoRegression.joblib']
```