

## KNN

Regression based on k-nearest neighbors.

The target is predicted by local interpolation of the targets associated of the nearest neighbors in the training set.

One thing to note KNN wont require any transformation so we are opting power transformation here!

```
In [ ]: # lets import some dependencies

import warnings
warnings.filterwarnings('ignore')

import matplotlib.pyplot as plt
%config InlineBackend.figure_format = 'retina'
plt.style.use('ggplot')
import pandas as pd
import numpy as np

import seaborn as sns
sns.set()

from sklearn.model_selection import train_test_split, KFold
from sklearn.model_selection import GridSearchCV, cross_val_score, learning_curve
from sklearn.model_selection import validation_curve

from sklearn.neighbors import KNeighborsRegressor

from sklearn.metrics import mean_absolute_error, r2_score

from joblib import load,dump
```

```
In [ ]: # loading our data
df = pd.read_csv("./Data/data.csv",sep=",")
df.drop(['Unnamed: 0'], axis=1, inplace=True) # There were some formatting issues
# writing the csv
```

```
In [ ]: df.head()
```

```
Out[ ]:   DISTRICT UPAZILA STATION_ID STATION_NAME DATE RAIN_FALL(mm) LATITUDE LONGITUDE
0  Bandarban    Lama     CL317      Lama  01-jan-2017        0.0    21.81      92.1
1  Bandarban    Lama     CL317      Lama  02-jan-2017        0.0    21.81      92.1
2  Bandarban    Lama     CL317      Lama  03-jan-2017        0.0    21.81      92.1
```

	DISTRICT	UPAZILA	STATION_ID	STATION_NAME	DATE	RAIN_FALL(mm)	LATITUDE	LONGITUDI
3	Bandarban	Lama	CL317	Lama	04-jan-2017	0.0	21.81	92.1
4	Bandarban	Lama	CL317	Lama	05-jan-2017	0.0	21.81	92.1

Defining our X and y

```
In [ ]: X = df['RAIN_FALL(mm)'].values.reshape(-1,1) # input feature
y = df['WATER_LEVEL(m)'].values.reshape(-1,1) # target feature
```

```
In [ ]: X.shape, y.shape
```

```
Out[ ]: ((1826, 1), (1826, 1))
```

Making the train test split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(
    X,y, test_size=0.2, random_state=17, shuffle=True
)
```

## Model Building

Initialize our CV

```
In [ ]: kflod = KFold(n_splits=5,shuffle=True, random_state=17)
```

Initialize the KNN model with default parameters

- n\_neighbors: 5
- weights: uniform
- algorithm: auto
- leaf\_size: 30
- p: 2 (power parameter of Minkowski metric)
- metric: minkowski
- metric\_params: None

```
In [ ]: knn = KNeighborsRegressor(n_jobs=-1)
```

As usual, lets first check our CV scores with the default parameters!

```
In [ ]: results = cross_val_score(
    knn,
    X_train,
    y_train,
    cv=kflod,
```

```
        scoring='neg_mean_absolute_error'
    )
results.mean()
```

```
Out[ ]: 0.4486728767123287
```

Note:

- This is better than our linear models!

Checking for train-test accuracy

```
In [ ]: knn.fit(X_train,y_train)
```

```
Out[ ]: KNeighborsRegressor(n_jobs=-1)
```

```
In [ ]: # accuracy on the train set
knn_pred = knn.predict(X_train)
mean_absolute_error(y_train,knn_pred)
```

```
Out[ ]: 0.3863758904109589
```

```
In [ ]: # accuracy on the test set
knn_pred = knn.predict(X_test)
mean_absolute_error(y_test,knn_pred)
```

```
Out[ ]: 0.4345224043715846
```

## Learning Curve


Lets plot the learning curve to see how our dataset's size influence the model accuracy

```
In [ ]: # Helper function
alphas = np.logspace(-2, 0, 20)
def plot_with_err(x, data, **kwargs):
    mu, std = data.mean(1), data.std(1)
    lines = plt.plot(x, mu, "-", **kwargs)
    plt.fill_between(
        x,
        mu - std,
        mu + std,
        edgecolor="none",
        facecolor=lines[0].get_color(),
        alpha=0.2,
    )
```

```
In [ ]: def plot_learning_curve():
    train_sizes = np.linspace(0.05,1,20)
    N_train, val_train, val_test = learning_curve(
        knn,X,y, train_sizes=train_sizes, cv=kfold, scoring='neg_mean_absolute_error')
    plot_with_err(N_train, val_train, label="training scores")
    plot_with_err(N_train, val_test, label="validation scores")
```

```
plt.xlabel("Training Set Size")
plt.ylabel("Negative MAE")
plt.legend()
plt.grid(True);
```

In [ ]: `plot_learning_curve()`



Note:

- From this we conclude, adding more data to this stage, wont help us to improve the model!

## Validation Curves

```
In [ ]: # helper function
def plot_with_err(x, data, **kwargs):
    mu, std = data.mean(1), data.std(1)
    lines = plt.plot(x, mu, "-", **kwargs)
    plt.fill_between(
        x,
        mu - std,
        mu + std,
        edgecolor="none",
        facecolor=lines[0].get_color(),
        alpha=0.2,
    )
```

```
In [ ]: def plot_validation_curve(param_grid,param,estimator):
    val_train, val_test = validation_curve(
        estimator=estimator,
        X=X_train,
        y=y_train,
        param_name=param,
        param_range=param_grid,
        cv=kfold,
        scoring="neg_mean_absolute_error",
        n_jobs=-1
    )

    plot_with_err(param_grid, val_train, label="training scores")
```


```

plot_with_err(param_grid, val_test, label="validation scores")
plt.xlabel(param)
plt.ylabel("Negative MAE")
plt.legend()
plt.grid(True);

```

- Number of Neighbors

In [ ]: `plot_validation_curve([5, 10, 15, 20, 30, 50, 75, 100], 'n_neighbors', knn)`



from figure its clear that the optimal range for number of neighbors is between 40 to 60

- weights


In [ ]: `knn = KNeighborsRegressor(n_neighbors=45, n_jobs=-1)  
plot_validation_curve(['uniform', 'distance'], 'weights', knn)`



There will be higher variance in the model if we use distance as our weights!

- algorithm


```
In [ ]: plot_validation_curve(['ball_tree', 'kd_tree', 'brute'], 'algorithm', knn)
```



there's no significant difference in algorithms

- leaf size


```
In [ ]: knn = KNeighborsRegressor(n_neighbors=45, algorithm='kd_tree', n_jobs=-1)
plot_validation_curve([5, 10, 15, 20, 30, 40, 50, 60], 'leaf_size', knn)
```



No significant difference here too

- p

```
In [ ]: plot_validation_curve([1, 2], 'p', knn)
```



doesn't matter which norm we use!

## Hyper Parameter Tuning

Based on the knowledge of the validation curves we plotted lets do the hyperparameter tuning of KNN in order to find the best model!

```
In [ ]: knn_params = {
    'n_neighbors': [40, 45, 50, 55, 60],
    'algorithm': ['ball_tree', 'kd_tree'],
    'weights': ['uniform', 'distacne']
}
knn = KNeighborsRegressor(n_jobs=-1)
gcv = GridSearchCV(knn, knn_params, n_jobs=-1, cv=kfold, verbose=1,
scoring='neg_mean_absolute_error')
```

```
In [ ]: # training the model
gcv.fit(X_train, y_train)
```

```
Out[ ]: Fitting 5 folds for each of 20 candidates, totalling 100 fits
GridSearchCV(cv=KFold(n_splits=5, random_state=17, shuffle=True),
estimator=KNeighborsRegressor(n_jobs=-1), n_jobs=-1,
param_grid={'algorithm': ['ball_tree', 'kd_tree'],
'n_neighbors': [40, 45, 50, 55, 60],
'weights': ['uniform', 'distacne']},
scoring='neg_mean_absolute_error', verbose=1)
```

```
In [ ]: # the best score in CV
gcv.best_score_
```

```
Out[ ]: -0.4377368493150685
```

```
In [ ]: # the best estimator
gcv.best_estimator_
```

```
Out[ ]: KNeighborsRegressor(algorithm='ball_tree', n_jobs=-1, n_neighbors=50)
```

```
In [ ]: # the best parameters
gcv.best_params_
```

```
Out[ ]: {'algorithm': 'ball_tree', 'n_neighbors': 50, 'weights': 'uniform'}
```

Lets check the train-test accuracy of our best estimator

```
In [ ]: # accuracy on train set
knn_pred = gcv.best_estimator_.predict(X_train)
mean_absolute_error(y_train,knn_pred)
```

```
Out[ ]: 0.42500967123287675
```

```
In [ ]: # accuracy on the test set
knn_pred = gcv.best_estimator_.predict(X_test)
mean_absolute_error(y_test,knn_pred)
```

```
Out[ ]: 0.422436174863388
```


### Final Note

- Theres little to none overfitting occurring here
- we have significantly increase the accuracy of our model to .42 meters!

### Learning Curve

```
In [ ]: def plot_learning_curve():
    train_sizes = np.linspace(0.05,1,20)
    N_train, val_train, val_test = learning_curve(
        gcv.best_estimator_,X_train,y_train,
        train_sizes=train_sizes,cv=kfold,scoring='neg_mean_absolute_error'
    )
    plot_with_err(N_train, val_train, label="training scores")
    plot_with_err(N_train, val_test, label="validation scores")
    plt.xlabel("Training Set Size")
    plt.ylabel("Negative MAE")
    plt.legend()
    plt.grid(True);
```

```
In [ ]: plot_learning_curve()
```



## Saving the model

```
In [ ]: dump(gcv.best_estimator_, './Saved Model/knn.joblib')
```

```
Out[ ]: ['./Saved Model/knn.joblib']
```