

## SVR

Epsilon-Support Vector Regression.

The free parameters in the model are C and epsilon.

```
In [ ]: # lets import the dependencies

import warnings
warnings.filterwarnings('ignore')

import matplotlib.pyplot as plt
%config InlineBackend.figure_format = 'retina'
plt.style.use('ggplot')
import pandas as pd
import numpy as np

import seaborn as sns
sns.set()

from sklearn.model_selection import train_test_split, KFold
from sklearn.model_selection import GridSearchCV, cross_val_score, learning_curve
from sklearn.model_selection import validation_curve

from sklearn.preprocessing import PowerTransformer

from sklearn.svm import SVR

from sklearn.metrics import mean_absolute_error, r2_score
from joblib import load,dump
```

```
In [ ]: # loading the data
# loading our data
df = pd.read_csv("./Data/data.csv",sep=",")
df.drop(['Unnamed: 0'], axis=1, inplace=True) # There were some formatting issues
                                              # writing the csv
```

```
In [ ]: df.head()
```

```
Out[ ]: DISTRICT UPAZILA STATION_ID STATION_NAME DATE RAIN_FALL(mm) LATITUDE LONGITUDE
0 Bandarban Lama CL317 Lama 01-jan-2017 0.0 21.81 92.1
1 Bandarban Lama CL317 Lama 02-jan-2017 0.0 21.81 92.1
2 Bandarban Lama CL317 Lama 03-jan-2017 0.0 21.81 92.1
3 Bandarban Lama CL317 Lama 04-jan-2017 0.0 21.81 92.1
```

	DISTRICT	UPAZILA	STATION_ID	STATION_NAME	DATE	RAIN_FALL(mm)	LATITUDE	LONGITUDI
4	Bandarban	Lama	CL317	Lama	05-jan-2017	0.0	21.81	92.1

Defining our X and y

```
In [ ]: X = df['RAIN_FALL(mm)'].values.reshape(-1,1) # input feature
y = df['WATER_LEVEL(m)'].values.reshape(-1,1) # target feature
```

making the train test split

```
In [ ]: X_train, X_test, y_train, y_test = train_test_split(
    X,y, test_size=0.2, random_state=17, shuffle=True
)
```

## PreProcessing

### Power Transformation

```
In [ ]: pt = PowerTransformer(method='box-cox')
X_train_transformed = pt.fit_transform(X_train+0.000001)
X_test_transformed = pt.transform(X_test+0.000001)
```

```
In [ ]: pt.lambdas_
```

```
Out[ ]: array([-0.10662485])
```

```
In [ ]: y_train_transformed = pt.fit_transform(y_train+0.000001)
y_test_transformed = pt.transform(y_test+0.000001)
y_train_transformed = y_train_transformed.ravel()
y_test_transformed = y_test_transformed.ravel()
```

```
In [ ]: pt.lambdas_
```

```
Out[ ]: array([-5.34754941])
```

## Model Building

Initialize our CV

```
In [ ]: kfold = KFold(n_splits=5,shuffle=True,random_state=17)
```

Initialize Support Vector Regressor with default parameters:

- kernel: rbf
- tol : 1e-3

- C: 1
- epsilon: 0.1
- verbose:True
- max\_iteration: -1

```
In [ ]: svr = SVR()
```

CV Scores:

```
In [ ]: results = cross_val_score(
    svr,
    X_train_transformed,
    y_train_transformed,
    cv=kfold,
    scoring='neg_mean_absolute_error'
)
results.mean()
```

```
Out[ ]: 0.5760287054297919
```

Note:

- we got 0.44 in KNN

Train-test accuracy

```
In [ ]: svr.fit(X_train_transformed,y_train_transformed)
```

```
Out[ ]: SVR()
```

```
In [ ]: # accuracy on the train set
pred = svr.predict(X_train_transformed)
mean_absolute_error(y_train_transformed,pred)
```

```
Out[ ]: 0.5751854260156385
```

```
In [ ]: # accuracy on the test set
pred = svr.predict(X_test_transformed)
mean_absolute_error(y_test_transformed,pred)
```

```
Out[ ]: 0.570577574929561
```

## Learning Curve

lets plot the learning curve to see how our dataset's size influence the model accuracy

```
In [ ]: # Helper function
alphas = np.logspace(-2, 0, 20)
def plot_with_err(x, data, **kwargs):
    mu, std = data.mean(1), data.std(1)
    lines = plt.plot(x, mu, "-", **kwargs)
```

```

    plt.fill_between(
        x,
        mu - std,
        mu + std,
        edgecolor="none",
        facecolor=lines[0].get_color(),
        alpha=0.2,
    )
)

```

In [ ]:


```

def plot_learning_curve():
    train_sizes = np.linspace(0.05, 1, 20)
    N_train, val_train, val_test = learning_curve(
        svr, X_train_transformed, y_train_transformed, \
        train_sizes=train_sizes, cv=kfold, scoring='neg_mean_absolute_error'
    )
    plot_with_err(N_train, val_train, label="training scores")
    plot_with_err(N_train, val_test, label="validation scores")
    plt.xlabel("Training Set Size")
    plt.ylabel("Negative MAE")
    plt.legend()
    plt.grid(True);

```

In [ ]:

```
plot_learning_curve()
```



Note:

- its clear that training and validation curve converges, assuring no variance problem in the data

## Validation Curves

In [ ]:

```

# helper function
def plot_with_err(x, data, **kwargs):
    mu, std = data.mean(1), data.std(1)
    lines = plt.plot(x, mu, "-", **kwargs)
    plt.fill_between(
        x,
        mu - std,
        mu + std,
        edgecolor="none",

```


```
        facecolor=lines[0].get_color(),
        alpha=0.2,
    )
```

```
In [ ]: def plot_validation_curve(param_grid,param,estimator):
    val_train, val_test = validation_curve(
        estimator=estimator,
        X=X_train_transformed,
        y=y_train_transformed,
        param_name=param,
        param_range=param_grid,
        cv=kfold,
        scoring="neg_mean_absolute_error",
        n_jobs=-1
    )

    plot_with_err(param_grid, val_train, label="training scores")
    plot_with_err(param_grid, val_test, label="validation scores")
    plt.xlabel(param)
    plt.ylabel("Negative MAE")
    plt.legend()
    plt.grid(True);
```

- **C:** Regularization parameter. The strength of the regularization is inversely proportional to C.  
Must be strictly positive. The penalty is a squared L2 penalty.

```
In [ ]: plot_validation_curve(np.linspace(1,200,20), 'C', svr)
```



Note:


- After 75, there's no significant improvement
- **epsilon:** Epsilon in the epsilon-SVR model. It specifies the epsilon-tube within which no penalty is associated in the training loss function with points predicted within a distance epsilon from the actual value.

```
In [ ]:
```

```

svr = SVR(C=80)
plot_validation_curve(np.linspace(0.1,10,10), 'epsilon', svr)

```



## Hyper-Parameter Tuning:

```

In [ ]:
svr_params = {
    'C': np.linspace(1,200,20),
    'epsilon': np.linspace(0.01,10,10)
}
svr = SVR()
gcv = GridSearchCV(svr, svr_params, n_jobs=-1, cv=kfold, verbose=1,
                    scoring='neg_mean_absolute_error')

```

```

In [ ]:
# training the model
gcv.fit(X_train_transformed, y_train_transformed)

```

Fitting 5 folds for each of 200 candidates, totalling 1000 fits

```

Out[ ]:
GridSearchCV(cv=KFold(n_splits=5, random_state=17, shuffle=True),
             estimator=SVR(), n_jobs=-1,
             param_grid={'C': array([ 1.          ,  11.47368421,  21.94736842,
            32.42105263,
               42.89473684,  53.36842105,  63.84210526,  74.31578947,
               84.78947368,  95.26315789, 105.73684211, 116.21052632,
              126.68421053, 137.15789474, 147.63157895, 158.10526316,
              168.57894737, 179.05263158, 189.52631579, 200.        ]),
             'epsilon': array([ 0.01,  1.12,  2.23,  3.34,  4.45,
               5.56,  6.67,  7.78,  8.89,
              10.        ])},
             scoring='neg_mean_absolute_error', verbose=1)

```

```

In [ ]:
# the best score
gcv.best_score_

```

```

Out[ ]:
-0.5636756238996249

```

```

In [ ]:
# the best estimators parameters
gcv.best_params_

```

```
Out[ ]: {'C': 179.05263157894734, 'epsilon': 0.01}
```

```
In [ ]: # Lets check the train test accuracy of our model  
svr_pred = gcv.best_estimator_.predict(X_train_transformed)  
mean_absolute_error(y_train_transformed,svr_pred)
```

```
Out[ ]: 0.5627343535124258
```


```
In [ ]: # Lets check the test test accuracy of our model  
svr_pred = gcv.best_estimator_.predict(X_test_transformed)  
mean_absolute_error(y_test_transformed,svr_pred)
```

```
Out[ ]: 0.5623383600494728
```

## Learning Curve

```
In [ ]: def plot_learning_curve():  
    train_sizes = np.linspace(0.05,1,20)  
    N_train, val_train, val_test = learning_curve(  
        gcv.best_estimator_,X_train_transformed,y_train_transformed,  
        train_sizes=train_sizes,cv=kfold,scoring='neg_mean_absolute_error'  
    )  
    plot_with_err(N_train, val_train, label="training scores")  
    plot_with_err(N_train, val_test, label="validation scores")  
    plt.xlabel("Training Set Size")  
    plt.ylabel("Negative MAE")  
    plt.legend()  
    plt.grid(True);
```

```
In [ ]: plot_learning_curve()
```



## Saving Model

```
In [ ]: dump(gcv.best_estimator_, './Saved Model/svr.joblib')
```

```
Out[ ]: ['./Saved Model/svr.joblib']
```

